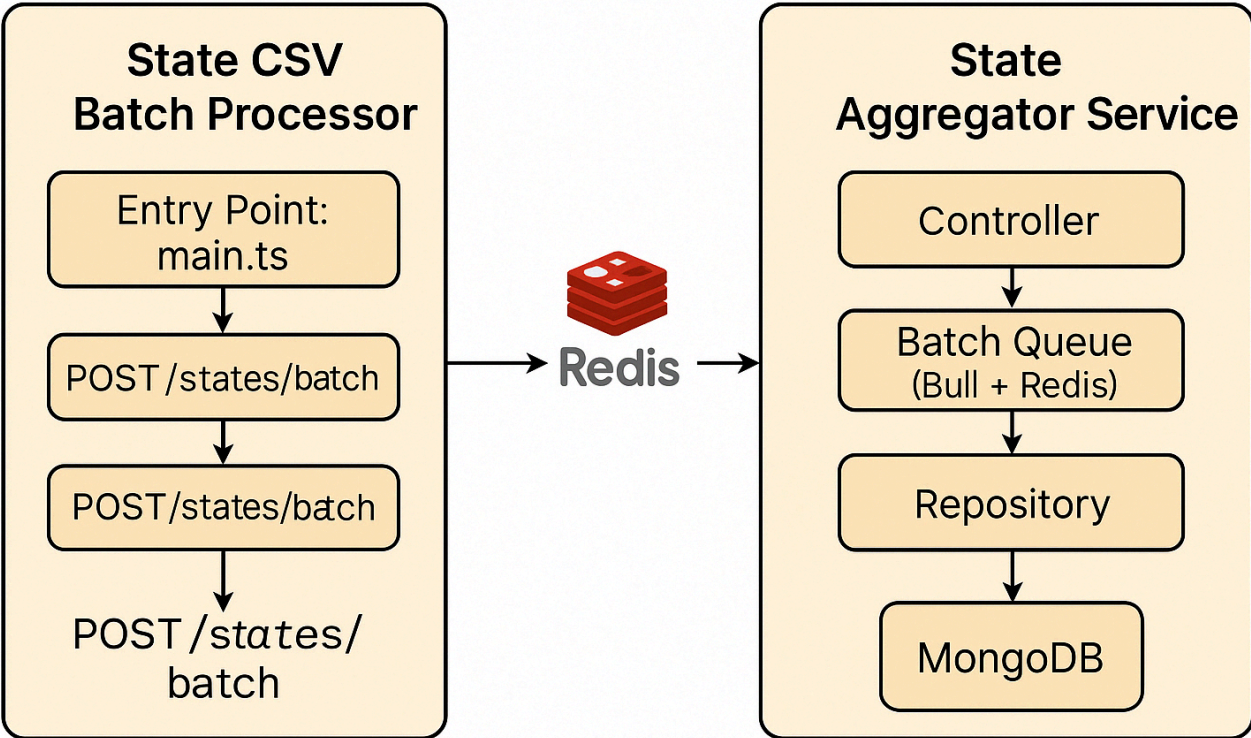


Architecture Design Documentation

Architecture Design (Diagram)



Architecture Design

Layered View

```
State CSV Batch Processor
-----
[Entry Point: main.ts]
|
[CsvReaderService]
|
[Batching & Validation]
|
[HTTP Client (axios)]
|
[POST /states/batch]
|
State Aggregator Service
-----
```

```
[Controller]
  |
[Service]
  |
[Batch Queue (Bull + Redis)]
  |
[Repository]
  |
[MongoDB]
```

Objective

The objective of this architecture is:

- **Automated Data Ingestion:** Efficiently process a large CSV file containing phone records, extracting and batching state population data.
- **Reliable Data Transfer:** Send this data in manageable batches to a backend service using HTTP, ensuring scalability and fault tolerance.
- **Centralized Aggregation & Storage:** The backend (State Aggregator Service) receives these batches, aggregates the population by state, and persists the results in a MongoDB database.
- **Asynchronous & Scalable Processing:** Use Redis as a queue backend (via Bull) to decouple HTTP request handling from heavy batch processing, enabling horizontal scaling and reliability.
- **Easy Data Access:** Expose REST API endpoints so users or other systems can query the aggregated state population data, monitor ingestion, and manage records.

Redis in the Architecture

- **Purpose:** Redis is used as a fast, in-memory data store to back the Bull queue system. This enables the State Aggregator Service to queue incoming batch jobs and process them asynchronously.
- **How it works:** When a batch is received via the `/states/batch` endpoint, it is placed into a Redis-backed queue. Worker processes (which can be scaled horizontally) pull jobs from this queue and perform aggregation and storage in MongoDB.
- **Benefits:**
 - Decouples HTTP requests from processing, improving API responsiveness.
 - Supports retry, delayed, and scheduled jobs.
 - Enables monitoring of job status and queue health.
 - Facilitates horizontal scaling—multiple workers can process jobs from the same queue.

In summary:

- State CSV Batch Processor automates reading and batching of CSV data, sending it to the backend.
- State Aggregator Service receives, queues (via Redis), aggregates, and stores the data, providing APIs for access and management.
- Redis (with Bull) acts as a reliable, scalable queue for batch processing, ensuring that heavy workloads do not block API requests and can be processed efficiently in the background.
- MongoDB stores the final, aggregated state population data.

This separation ensures each part of the system is focused, maintainable, scalable, and robust for production workloads.